

Preface

Many compiler books have been published over the years, so why another one? Let me be perfectly clear. This book is designed primarily for use as a textbook in a one-semester course for undergraduate students and beginning graduate students. The only prerequisites for this book are familiarity with basic algorithms and data structures (lists, maps, recursion, etc.), a rudimentary knowledge of computer architecture and assembly language, and some experience with the Kotlin programming language or a closely related language such as Java or C#. (A separate edition of this book is available that uses Java as the implementation language for the compiler.) Most undergraduate computer science majors will have covered these topics in their first two years of study. Graduate students who have never had a course in compilers will also find the book useful, especially if they undertake some of the more challenging project exercises described in Appendix B.

A complete study of compilers could easily fill several graduate-level courses, and therefore some simplifications and compromises are necessary for a one-semester course that is accessible to undergraduate students. Following are some of the decisions made in order to accommodate the goals of this book.

1. The book has a narrow focus as a project-oriented course on compilers. Compiler theory is kept to a minimum, but the project orientation retains the “fun” part of studying compilers.
2. The source language being compiled is relatively simple, but it is powerful enough to be interesting and challenging. It has scalar data types, strings, arrays, records, procedures, functions, and parameters, but it relegates many other interesting language features to the project exercises. Most undergraduate students will find it challenging just to complete a compiler for the basic project language without any additional features. Graduate students will want to extend the basic project language with features outlined in the exercises.
3. The target language is assembly language for a virtual machine with a stack-based architecture, similar to but much simpler than the Java Virtual Machine (JVM). This approach greatly simplifies code generation. First, it eliminates the need to deal with general-purpose registers. And second, relative addresses for branch instructions are handled by the assembler, simplifying the amount of work that needs to be done by the compiler. Both an assembler and an emulator for the virtual machine are provided in the book’s GitHub repository.
4. No special compiler-related tools are required or used within the book. Students require access only to a Kotlin compiler and a text editor, but most students will want to use Kotlin with an Integrated Development Environment (IDE) such as IntelliJ IDEA or Eclipse. Compiler-related tools such as scanner generators or parser generators could simplify certain tasks involved in building a compiler, but I believe that the approach used in this book makes the structure of the compiler more transparent. Students who wish to use compiler-related tools are welcome to do so, but they will need to look elsewhere to learn how to use these tools. Examples of freely available compiler-related tools include ANTLR, Coco/R, Flex/Bison, Lex/Yacc,

JavaCC, and Eclipse Xtext. In addition, while the presentation of the book uses Kotlin, students are free to select an alternative implementation language. Languages that support recursion and object-oriented programming will work best with the approach used in this book. Examples include Java, C++, C#, Python, Scala, and Swift.

5. One very important component of a compiler is the parser, which verifies that a source program conforms to the language syntax and produces an intermediate representation of the program that is suitable for additional analysis and code generation. There are several different approaches to parsing, but in keeping with the emphasis on a one-semester course, this book covers only one approach – recursive descent parsing with several lookahead tokens. The previous edition of this book used only one lookahead token, but changes to the source language for this edition motivate the need for additional lookahead tokens. Plus, it is instructive to learn how to implement a parser using multiple lookahead tokens, especially since multiple lookahead tokens can be implemented simply and efficiently.
6. Missing from the book are a lot of favorite compiler topics such as finite automata, bottom-up parsing, attribute grammars, heap management, register allocation, and data-flow analysis. What remains fits nicely into a one-semester, project-oriented course for undergraduate students and beginning graduate students.

Why Study Compilers?

Relatively few people write commercial compilers, but I believe that the study of compiler design is an important part of a computer science education since compilers and closely related language translators are the primary tools for software development. Having a fundamental knowledge of how compilers actually work can improve one's ability to write good programs. In addition, learning how to write a compiler gives experience on working with a moderately large program, and it makes an excellent case study for software engineering and good program design. And the techniques used in compiler design are useful for other purposes such as writing a test harness, a support tool for process control, or an interpreter for a small, special-purpose language. One of my former students applied compiler technology to a commercial tool for testing software.

The Course Project

This book discusses implementation of a compiler for a relatively small programming language named CPRL (for Compiler Project Language), which was designed for teaching basics of compiler construction. The target language is assembly language for CVM (CPRL Virtual Machine), a simple stack-based virtual machine. Together we will build a compiler slowly, one step at a time, with lots of template Kotlin code in the book and even more in the book's GitHub repository to guide you through the process. Students are expected to download and study the Kotlin code from the book repository as part of their learning experience. The end result is likely the largest single program that most undergraduate students will encounter.

Appendix A describes the decomposition of the overall project of developing a CPRL compiler into 11 smaller subprojects (numbered 0-10). Students should complete the subprojects in the specified order since each one builds on the previous.

Also, we will build the parser in three separate subprojects, focusing initially on the syntax as defined by the context-free grammar, then adding error recovery, and finally constructing abstract syntax trees for subsequent phases of constraint analysis and initial code generation. In the real world we would likely combine all three steps into one, but the division into three separate subprojects has been shown to provide a better understanding of the concepts and overall process for students new to compiler design.

Appendix B describes a number of extensions or variations to the basic compiler project outlined in this book. Some are more challenging than others. Ambitious undergraduate students and most graduate students will want to attempt one or more of these exercises.

Changes in the Second Edition

The second edition retains the objectives, emphasis, and general structure of the first edition, but there are a number of significant changes that can be grouped into four broad categories.

1. CPRL Language Changes.

- The CPRL language is now a little more like C-based languages and less like Pascal-based languages. For example, reserved words “begin” and “end” have largely been replaced by braces “{” and “}”, and program execution now starts with a parameterless procedure named “main()”. Similarly, subroutine declarations and calls now require parentheses even if there are no parameters.
- Subprograms, including `main()`, can now be called before they have been declared, but initial declarations (i.e., variables, constants, and types) must still appear before being used.
- String types are now fully supported (the previous edition supported only string literals), and record types (a.k.a., structures) were added.
- Similar to C, array parameters in CPRL are now passed by reference whether or not they are declared as such, even for function parameters.

2. CVM Changes.

- There are new instructions to support string input and assignment plus several new optimized versions of existing instructions.
- Instructions that read from standard input (e.g., `GETINT`) now write directly to the variable address found on the top of the stack. Previously these instructions pushed the input value onto the stack, and a store instruction was required to update the variable.
- The `CMP` (compare) instruction was removed. So, for example, branch instructions like `BL` (branch less than) and `BGE` (branch greater or equal) now pop two integers

off the stack, compare them, and branch accordingly. This approach is similar to the way branch instructions work in the JVM and other stack-based architectures.

3. Compiler Implementation Changes.

- The use of null values has been greatly reduced, especially with respect to abstract syntax trees. For example, most parser methods now return either `EmptyExpression`, `EmptyInitialDecl`, `EmptySubprogramDecl`, or `EmptyStatement` rather than null when a syntax error is encountered.
- The parser can now use more than one lookahead token to make parsing decisions.

4. Other Changes.

- The notation for context-free grammars was changed to more closely resemble the one defined in ISO/IEC 14977, an international standard for EBNF notation. For example, “zero or more” is now represented by braces “{“ and “}”.
- A more complete discussion of scopes has been added together with a cleaner implementation of scopes in the compiler.

In addition, there are two new chapters, several new exercises, many new paragraphs, various rewordings to make the presentation easier to understand, and a few corrections. Overall I am extremely pleased with the final result.

Structure of the Book

The book chapters are organized in a natural progression from general concepts to step-by-step details for implementing the compiler project. Each chapter builds on the previous chapters, and forward references have been kept to a minimum. A few sections can be omitted or postponed without loss of continuity, but in general, the book is designed to be studied from front to back.

Chapter 1 introduces basic concepts and terminology in the study of compilers. Following the book by David A. Watt and Deryck F. Brown [Watt 2000], Chapter 1 uses tombstone diagrams as a visual aid to explain many of the concepts. Sections 1.2, 1.3, and 1.6 are required. The other sections can be omitted or postponed without impacting the material in subsequent chapters.

Chapter 2 presents an overview of the structure of a compiler. In a sense, it sets the context for the remaining chapters.

Chapter 3 provides an overview of how programming languages are defined, with detailed coverage of context-free grammars. The complete context-free grammar for CPRL, the source language for the compiler project, is defined in Appendix D.

Chapter 4 gives a brief overview of CPRL. A more complete definition of CPRL is provided in Appendix C.

Chapter 5 presents the details of implementing a scanner for CPRL, including an overview of several related classes found in the book’s GitHub repository.

Chapter 6 is the longest and one of the most complicated chapters in the book. It includes a discussion of grammar analysis to compute first sets and follow sets plus details on the implementation of a recursive descent parser for CPRL. As defined in the compiler project, the initial version of the parser concentrates primarily on verification that a source program conforms to the syntax defined by the context-free grammar for CPRL.

Chapter 7 extends the parser developed in Chapter 6 to perform error recovery so that multiple errors can be detected and reported.

Chapter 8 contains a detailed discussion of abstract syntax trees. Abstract syntax trees provide an intermediate representation for source programs that can be used for additional analysis and code generation. In this chapter, the parsing methods are modified to construct an abstract syntax tree for the source program.

Chapter 9 uses the abstract syntax trees from Chapter 8 to perform constraint analysis, with emphasis on checking conformance to CPRL's type rules as defined for CPRL/0, a major subset of CPRL.

Chapter 10 gives a brief overview of the CVM, the virtual machine that serves as the target for CPRL programs. The CVM has a stack architecture that is similar to but much simpler than the Java Virtual Machine (JVM). Appendix E contains additional information about the CVM including complete definitions for every CVM instruction.

Chapter 11 modifies the abstract syntax trees to generate assembly language for CPRL/0. The GitHub repository provides an assembler that can be used to generate the actual CVM machine code.

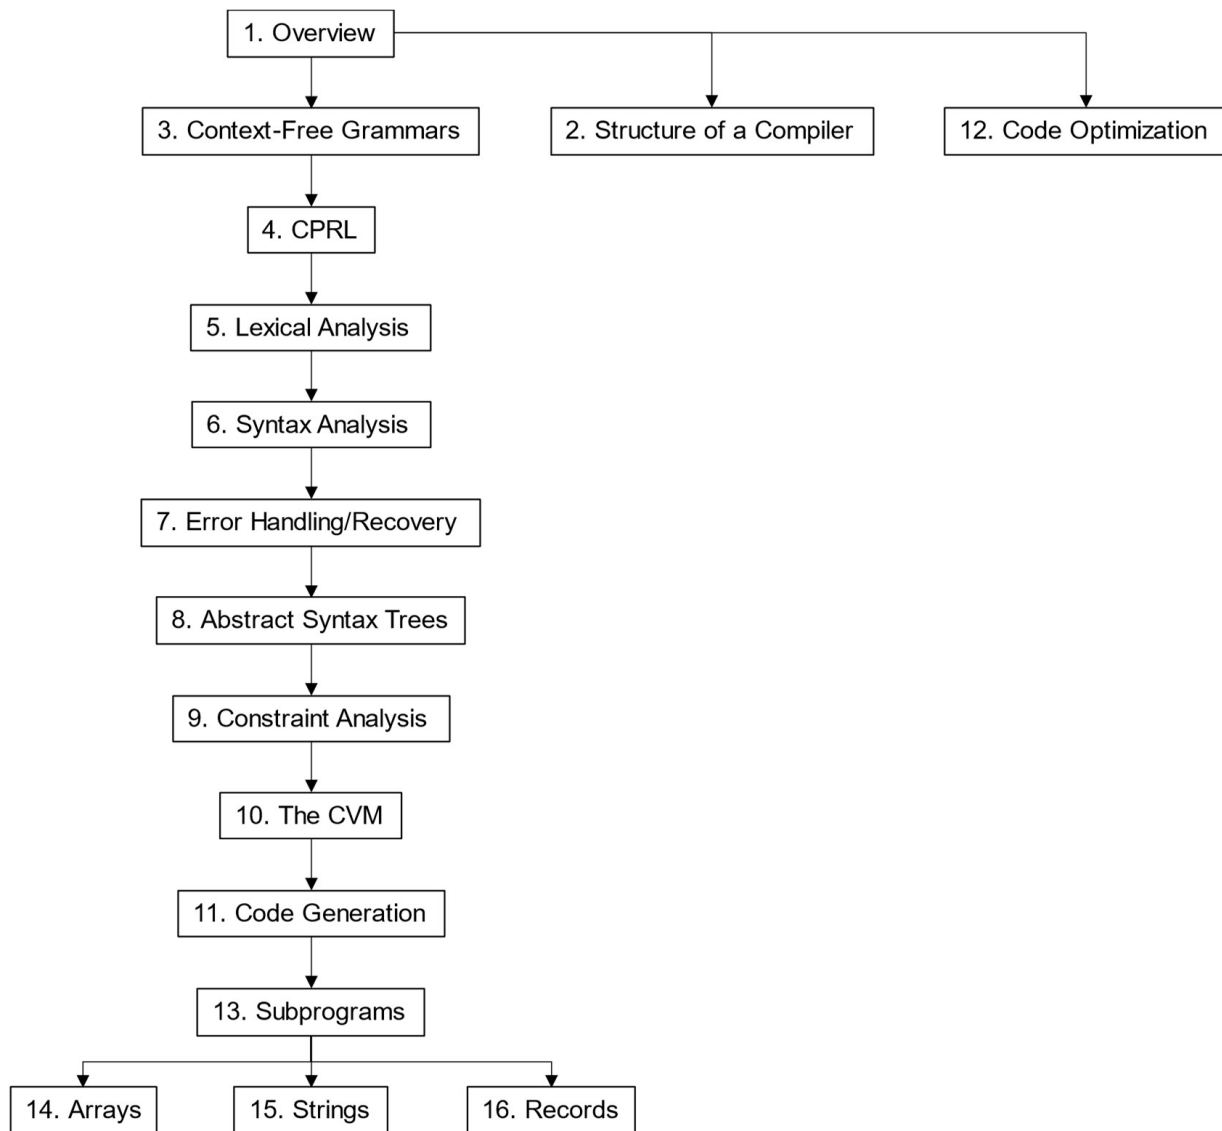
Chapter 12 presents an overview of code optimization. The assembler provided on the course web site implements several optimizations, but none of the basic compiler projects require mastery of the material from this chapter. Therefore, this chapter can be postponed if desired, but it should not be omitted.

Chapter 13 extends constraint analysis and code generation to handle subprograms. It includes additional coverage of scope and a detailed discussion of calling conventions and activation records. Similar to Chapter 6, this chapter is long and can be quite complicated for many undergraduate students. But mastery of this material is essential for a course on compilers, even if it means that less time is available for the last three chapters.

Chapters 14, 15, and 16 extend constraint analysis and code generation for arrays, strings, and records, respectively. Implementing records completes the basic compiler project, but students are always encouraged to attempt some of variations and extensions outlined in Appendix B.

Chapter Dependencies

As illustrated in the following diagram, Chapters 2 and 12 can be covered any time after Chapter 1, and Chapters 14-16 can be covered any time after Chapter 12. Otherwise, most chapters should be covered in the order presented in the book.



Book Resources

The repository for this book at <https://github.com/SoftMoore/CPRL-Kt-2nd> contains a number of related resources as follows.

- Kotlin source code that implements the CVM, the target machine for the compiler project.
- Kotlin source code that implements an assembler for the CVM. The compiler project targets assembly language for the CVM rather than the actual virtual machine.
- Kotlin source code for a disassembler; i.e., a program that takes CVM machine code and converts it back into assembly language. This program can be useful in trying to understand how machine code is laid out in memory.

- Kotlin source code or skeletal Kotlin source code for many of the classes described in the book so that students don't need to start from scratch to create their compilers. Much of the code for these classes reveal implementation ideas for other classes whose implementations are either not provided or are only partially provided. Students should begin each phase of their compiler project by trying to understand the related Kotlin source code that is provided.
- Many examples of correct and incorrect CPRL programs that can be used to test a compiler. The final compiler should reject all incorrect CPRL programs with appropriate error messages, and it should generate semantically equivalent assembly language for all correct CPRL programs. Using the provided assembler and CVM emulator, all correct CPRL programs can be run to compare observed and expected results. Students are strongly encouraged to develop additional test programs.
- Sample Windows command scripts and Bash shell scripts for running and testing various stages of the compiler. For example, rather than trying to compile and test programs one at a time, there are command/shell files for running the compiler on all CPRL source files in the current working directory. These command/shell files are useful for testing the compiler against collections of correct and incorrect CPRL programs.

Appendix A provides additional details about the resources available on the book repository and how those resources fit into the overall structure of the course compiler project.

Kotlin and Object-Orientation

As indicated in the book's subtitle, the language used herein for implementing a compiler is Kotlin, but other options are certainly possible. In fact, one of the project exercises listed in Appendix B is to write the compiler in a different language, and some of my former students have done this.

Since Kotlin has strong support for object-oriented programming, then most components of the compiler are structured in terms of classes and objects – the scanner is an object, the parser is an object, tokens are objects, etc. But the full power of “object-orientation” doesn't come into play until Chapter 8, where the details of abstract syntax trees are introduced. The approach used in this book for implementing abstract syntax trees fully exploits Kotlin's object-oriented facilities for inheritance and polymorphism. As an example, consider that code generation for a list of statements looks like the following.

```
for (stmt in statements)
    stmt.emit()
```

The list could include if statements, assignment statements, loop statements, etc., and therefore different statements generate code in different ways. But at the level of programming abstraction shown in this example, we only need to know that each statement has an `emit()` method that writes out the appropriate object code.

Similarly, the code to check constraints for a list of CPRL subprogram declarations looks like the following.

```
for (decl in subprogDecls)
    decl.checkConstraints()
```

Note that the subprogram declarations in the list could be either procedures or functions.

Which Version of Kotlin?

The exposition in this book and the resources provided in the book’s GitHub repository assume Kotlin/JVM (Kotlin targeted to the JVM platform) version 1.6 or later. For Java, version 14 or later is assumed. It is likely that early versions of Kotlin and/or Java can be used, but these are the recommended minimum versions for this book.

A Note from the Author About Formatting Source Code

Every programmer has his/her preferred way to format and indent source code. Should curly braces be aligned? Should code be indented 2, 3, or 4 spaces? What is the best way to format an `if` statement or a `while` loop? Sometimes arguments about how source code should be formatted seem to take on religious-like fervor. I avoid such arguments, for I have found the **one true way** that Kotlin source code should be formatted. ☺ But if you are a nonbeliever, please feel free to reformat the Kotlin source code according to your project standards or personal tastes. (FYI: The formatting style used in this book is loosely based on the GNU style, which is not very popular but which I have personally found to be the most readable. “Beauty is in the eye of the beholder.”)

Acknowledgements

When I was first learning about computers as a young man, I was in awe of compilers. That awe evolved into curiosity, which led to a lot of thought and study about how compilers are constructed, which led to a lot of contemplation about how best to teach an introductory course in compilers, which eventually led to my writing this book. Along the way I was guided and influenced by a number of teachers, colleagues, students, and books. I would like to acknowledge those that have been the most influential on my thinking.

I was first introduced to compilers many years ago when I “unofficially” audited a course on compiler design given by Richard LeBlanc at Georgia Institute of Technology (a.k.a., Georgia Tech). For reasons I can no longer remember, I was unable to take the course for credit, but auditing it, especially under Richard LeBlanc, was enough to motivate me to learn more. I was fortunate enough to have Richard LeBlanc later on for another course on programming language design.

My next introduction to compilers came in a professional development course given by Frank DeRemer and Tom Pennello, with a guest lecture or two by Bill McKeeman. They will not remember me, but I am grateful to have learned more about compilers from teachers and researchers of their calibers.

I have also been inspired by several compiler books, especially two of them that took a pedagogical approach similar to this one. The first book is *Brinch Hansen on Pascal Compilers* by Per Brinch Hansen (Prentice Hall, 1985). That book is a little out of date now, but it had one of the most readable descriptions of compilers when it was first released. A second, much more modern book is *Programming Language Processors in Java: Compilers and Interpreters* by David A. Watt and Deryck F. Brown (Prentice Hall 2000). I followed their treatment of tombstone diagrams when explaining compilers and interpreters in Chapter 1. Years ago I used the Brinch Hansen book as a textbook in my compiler courses, and more recently I used the Watt-Brown book several times when it was first published.

It is important to acknowledge former students at Johns Hopkins University and The Citadel and to thank these institutions for allowing me to explore my ideas about writing compilers. I am particularly grateful to the following students who had the most influence on my thinking about how to teach compilers: Rob Ring, Scott Stanchfield, Zack Aardahl, Ben Hunter, Gordon Finlay, Josh Terry, Davis Jeffords, Mike Dalpee, Matthew Blair, Mafer Contreras, Jared Johnson, Robert Roser, and Shiloh Smiles. Some of the earlier students at Johns Hopkins suffered through my courses as I was trying to crystallize my approach to teaching compilers, and I am grateful for their feedback. They would barely recognize my course if they took it today.

I must also acknowledge Vince Sigillito, who served as Chair of the Computer Science program at the Johns Hopkins Part-Time Programs in Engineering and Applied Science, for first allowing me to teach a course on compilers many years ago. I remember telling Vince that I wanted to teach a course on compilers so that I could learn more about them. He had no objections and even indicated that he had done something similar in the past. Compiler Design has evolved into my favorite course.

I would especially like to acknowledge the following individuals who provided invaluable advice and feedback on the presentation and exposition in this book: Richard LeBlanc, Art Pyster, Shankar Banik, and George Rudolph. I am most appreciative and humble that they spent part of their valuable time to assist me with this effort.

Finally, I would like to acknowledge Kayran Cox Moore, my wife of many years, for proofreading and providing invaluable feedback on several drafts of this book. She might not understand compilers or Kotlin programming, but she has a deep understanding of English grammar and sentence structure, and she has no reservations about correcting my errors or improving my writing. Any grammatical errors remaining in this book are a result of my stubborn refusal to follow her advice, or they were introduced after she proofread that part of the book. I also want to thank Kayran for being my anchor in life and the source for most of what is good about myself.

